



ALMA MATER STUDIORUM  
UNIVERSITÀ DI BOLOGNA  
DEPARTMENT OF COMPUTER SCIENCE AND ENGINEERING

# Real-Time Systems for Automation M

## 2. System Structures

# Notice

---

The course material includes slides downloaded from:

<http://codex.cs.yale.edu/avi/os-book/>

*(slides by Silberschatz, Galvin, and Gagne, associated with  
Operating System Concepts, 9th Edition, Wiley, 2013)*

and

<http://retis.sssup.it/~giorgio/rts-MECS.html>

*(slides by Buttazzo, associated with Hard Real-Time Computing  
Systems, 3rd Edition, Springer, 2011)*

which have been edited to suit the needs of this course.

The slides are authorized for personal use only.

Any other use, redistribution, and any for profit sale of the slides (in any form) requires the consent of the copyright owners.

# Previously

---

- ❑ The **operating system** act as **intermediary** between applications and hardware
- ❑ User applications can request the **services** of the OS through **system calls**
- ❑ **Program**: passive entity
- ❑ **Process**: program in execution, active entity, unit of work
  - ❑ Each process has a Process ID associated: **PID**



# Chapter 2: Operating-System Structures

---

- Operating System Services
- User and Operating System Interface
- System Calls
- Types of System Calls
- System Programs
- Operating-System Structure
- System Boot





# Operating System Services

- Operating systems provide an environment for execution of programs and services to programs and users
- One set of operating-system **services** provides functions that are helpful to the user:
  - **User interface** - Almost all operating systems have a user interface (**UI**).
    - ▶ Varies between **Command-Line (CLI)**, **Graphics User Interface (GUI)**, **Batch**
  - **Program execution** - The system must be able to load a program into memory and to run that program, and end its execution, either normally or abnormally (indicating error)
  - **I/O operations** - A running program may require I/O. For efficiency and protection, users cannot control I/O devices directly. OS must provide a means to do I/O *→ it needs to be regulated*
  - **File-system manipulation** - Programs need to read and write files and directories, create and delete them, search them, list file information, permission management. The OS must offer these functions





# Operating System Services

- **Communications** – Processes may exchange information on the same computer or between computers over a network

- ▶ Communications may be via either:

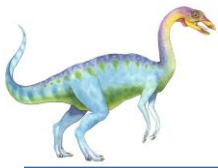
Same  
"PID" ←  
send/  
receive

- shared memory (in logical address space of the processes)
- message passing (in the kernel, packets moved by the OS which offers send/receive operations to the processes)

- **Error detection** – OS needs to be constantly aware of errors

- ▶ May occur in the CPU and memory hardware, in I/O devices, in user program
- ▶ For each type of error, OS should take the appropriate action to ensure correct and consistent computing
- ▶ Debugging facilities: OS provides the environment for a controlled, "interruptable" execution for testing; or core/memory dump facilities for post-mortem analysis.





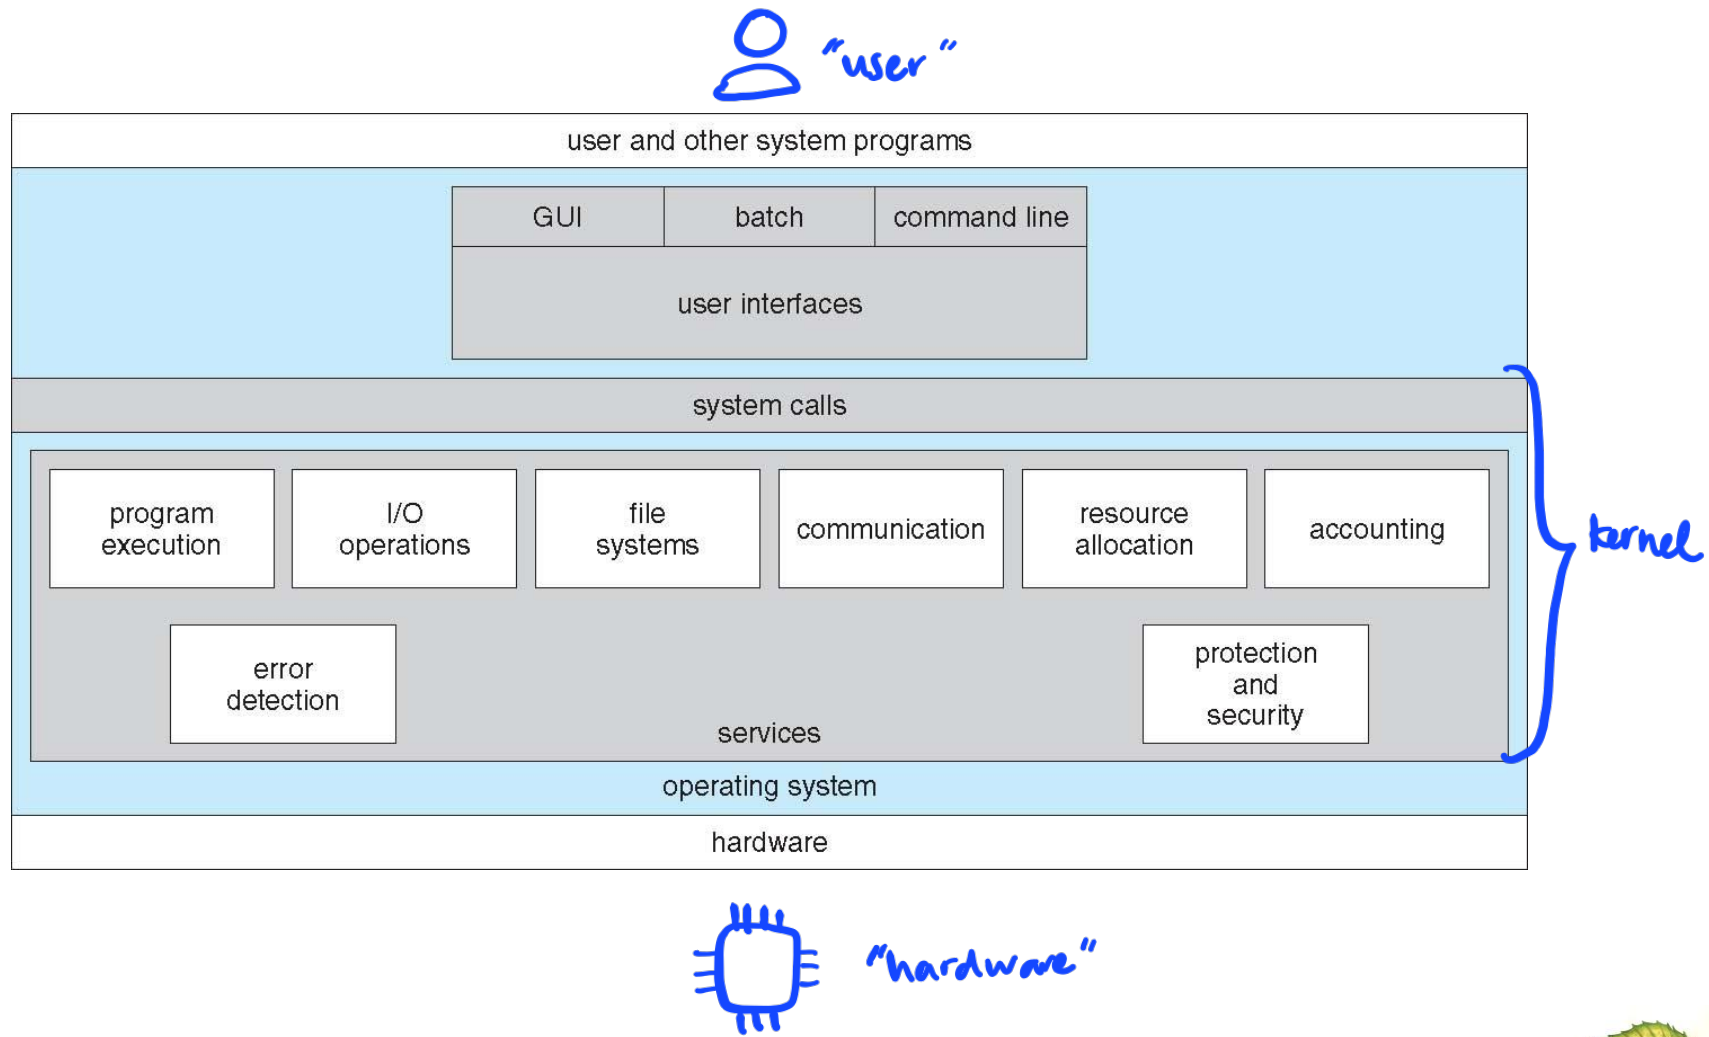
# Operating System Services

- Another set of OS functions exists for ensuring the **efficient operation** of the system itself via resource sharing
  - **Resource allocation** - When **multiple users or multiple jobs running concurrently**, resources must be allocated to each of them
    - ▶ Many types of resources - Some (such as CPU cycles, main memory, and file storage) may have special allocation code, others (such as I/O devices) may have general request and release code
  - **Logging** - To keep track of which users use how much and what kinds of computer resources (for accounting or statistics)
  - **Protection and security** - The owners of information stored in a multiuser or networked computer system may want to control use of that information, concurrent processes should not interfere with each other
    - ▶ **Protection** involves ensuring that all access to system resources is controlled
    - ▶ **Security** of the system from outsiders requires user authentication, extends to defending external I/O devices from invalid access attempts
    - ▶ If a system is to be protected and secure, precautions must be instituted throughout it. A chain is only as strong as its weakest link.





# A View of Operating System Services







# User Operating System Interface - CLI

- CLI or **command interpreter** allows **direct command entry**
  - it makes the user easier to call system calls
- Sometimes implemented in kernel, sometimes by systems program
- Sometimes multiple flavors implemented – **shells**
  - different ways to interact
- Primarily fetches a command from user, interprets and executes it
  - type command
  - OS read
  - ▶ built-in commands
  - ▶ external executions: in foreground (default) or background (using symbol & )
    - usually comes from the library





# Bourne Shell Command Interpreter

who

i/o stats

Ctrl+c

list

```
Default
New Info Close Execute Bookmarks

PBG-Mac-Pro:~ pbq$ w
15:24 up 56 mins, 2 users, load averages: 1.51 1.53 1.65
USER      TTY      FROM            LOGIN@   IDLE   WHAT
pbq       console  -                14:34    50    -
pbq       s000    -                15:05     -    w

PBG-Mac-Pro:~ pbq$ iostat 5

            disk0            disk1            disk10            cpu            load average
      KB/t tps MB/s      KB/t tps MB/s      KB/t tps MB/s      us sy id 1m 5m 15m
    33.75 343 11.30    64.31 14 0.88    39.67 0 0.02    11 5 84 1.51 1.53 1.65
     5.27 320 1.65      0.00 0 0.00      0.00 0 0.00     4 2 94 1.39 1.51 1.65
     4.28 329 1.37      0.00 0 0.00      0.00 0 0.00     5 3 92 1.44 1.51 1.65

^C
PBG-Mac-Pro:~ pbq$ ls
Applications          Music                  WebEx
Applications (Parallels)  Pando Packages       config.log
Desktop               Pictures              getsmartdata.txt
Documents              Public                imp
Downloads              Sites                 log
Dropbox                Thumbs.db             panda-dist
Library                Virtual Machines      prob.txt
Movies                 Volumes               scripts

PBG-Mac-Pro:~ pbq$ pwd
/Users/pbq

PBG-Mac-Pro:~ pbq$ ping 192.168.1.1
PING 192.168.1.1 (192.168.1.1): 56 data bytes
64 bytes from 192.168.1.1: icmp_seq=0 ttl=64 time=2.257 ms
64 bytes from 192.168.1.1: icmp_seq=1 ttl=64 time=1.262 ms

^C
--- 192.168.1.1 ping statistics ---
2 packets transmitted, 2 packets received, 0.0% packet loss
round-trip min/avg/max/stddev = 1.262/1.760/2.257/0.498 ms
PBG-Mac-Pro:~ pbq$
```

print working directory





# Basic Linux CLI

print smth

process status

how the user interacts that issues traps  
password concatenation

MANUAL

- Sample commands: echo, ls, cd, ps, pwd, cat, less, **man**
- File system: /bin, /usr/bin, /mount, /usr, /home, ...
- Commands pipeline: ls /usr/bin | less
- File creation: `> file1`
- Output redirect: `ls /usr/bin > output` (overwrite); `ls /usr/bin >> output` (append)
- Background execution: `sleep 5 &`

show less message

write smth

create an empty file

the command doesn't do anything until the process finished

read the output and overwrite the file





# System Calls instructions

---

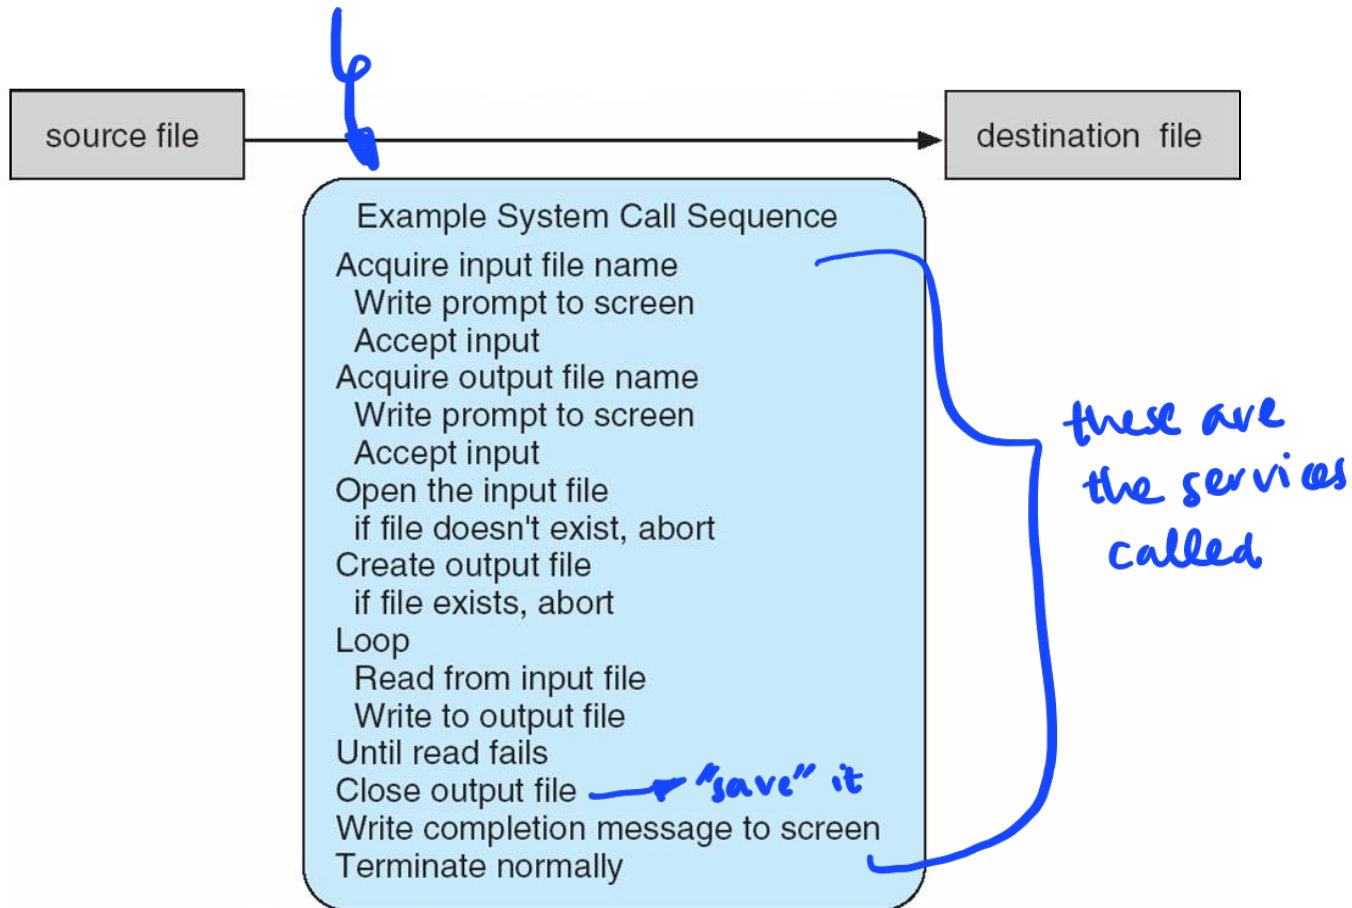
- Programming interface to the services provided by the OS
- Example: *copy the contents of one file to another file*





# Example of System Calls

- System call sequence to copy the contents of one file to another file





# System Calls

- Programming interface to the services provided by the OS
- Typically written in a high-level language (C or C++). Some low-level tasks in assembly-language
- Mostly accessed by programs via a high-level **Application Program Interface (API)** rather than direct system call use
  - E.g., in (most) Linux distro, you call `open(...)`, not `sys_open(...)`
- Three most common APIs are:
  - Win32 API for Windows,
  - POSIX API for POSIX-based systems (including virtually all versions of UNIX, Linux, and Mac OS X)
  - Java API for the Java virtual machine (JVM)
- Why use **APIs** rather than system calls?
  - portability ✓
  - simplicity ✓

*how the service is written*

*to avoid direct system calls use*





# Example of Standard API

## EXAMPLE OF STANDARD API

As an example of a standard API, consider the `read()` function that is available in UNIX and Linux systems. The API for this function is obtained from the man page by invoking the command

`man read`

on the command line. A description of this API appears below:

```
#include <unistd.h>
ssize_t read(int fd, void *buf, size_t count)
```

Diagram illustrating the components of the `read()` function signature:

return value	function name	parameters
<code>ssize_t</code>	<code>read</code>	<code>(int fd, void *buf, size_t count)</code>

Handwritten annotations:

- A blue arrow points from the word "variable" to `ssize_t`.
- A blue arrow points from the word "type" to `size_t`.
- A blue arrow points from the word "returns" to the return value `ssize_t`.
- A blue arrow points from the word "or" to the return value `ssize_t`.
- A blue arrow points from the word "-1" to the return value `ssize_t`.

A program that uses the `read()` function must include the `unistd.h` header file, as this file defines the `ssize_t` and `size_t` data types (among other things). The parameters passed to `read()` are as follows:

- `int fd`—the file descriptor to be read
- `void *buf`—a buffer where the data will be read into
- `size_t count`—the maximum number of bytes to be read into the buffer

On a successful read, the number of bytes read is returned. A return value of 0 indicates end of file. If an error occurs, `read()` returns `-1`.





# System Call Implementation

---

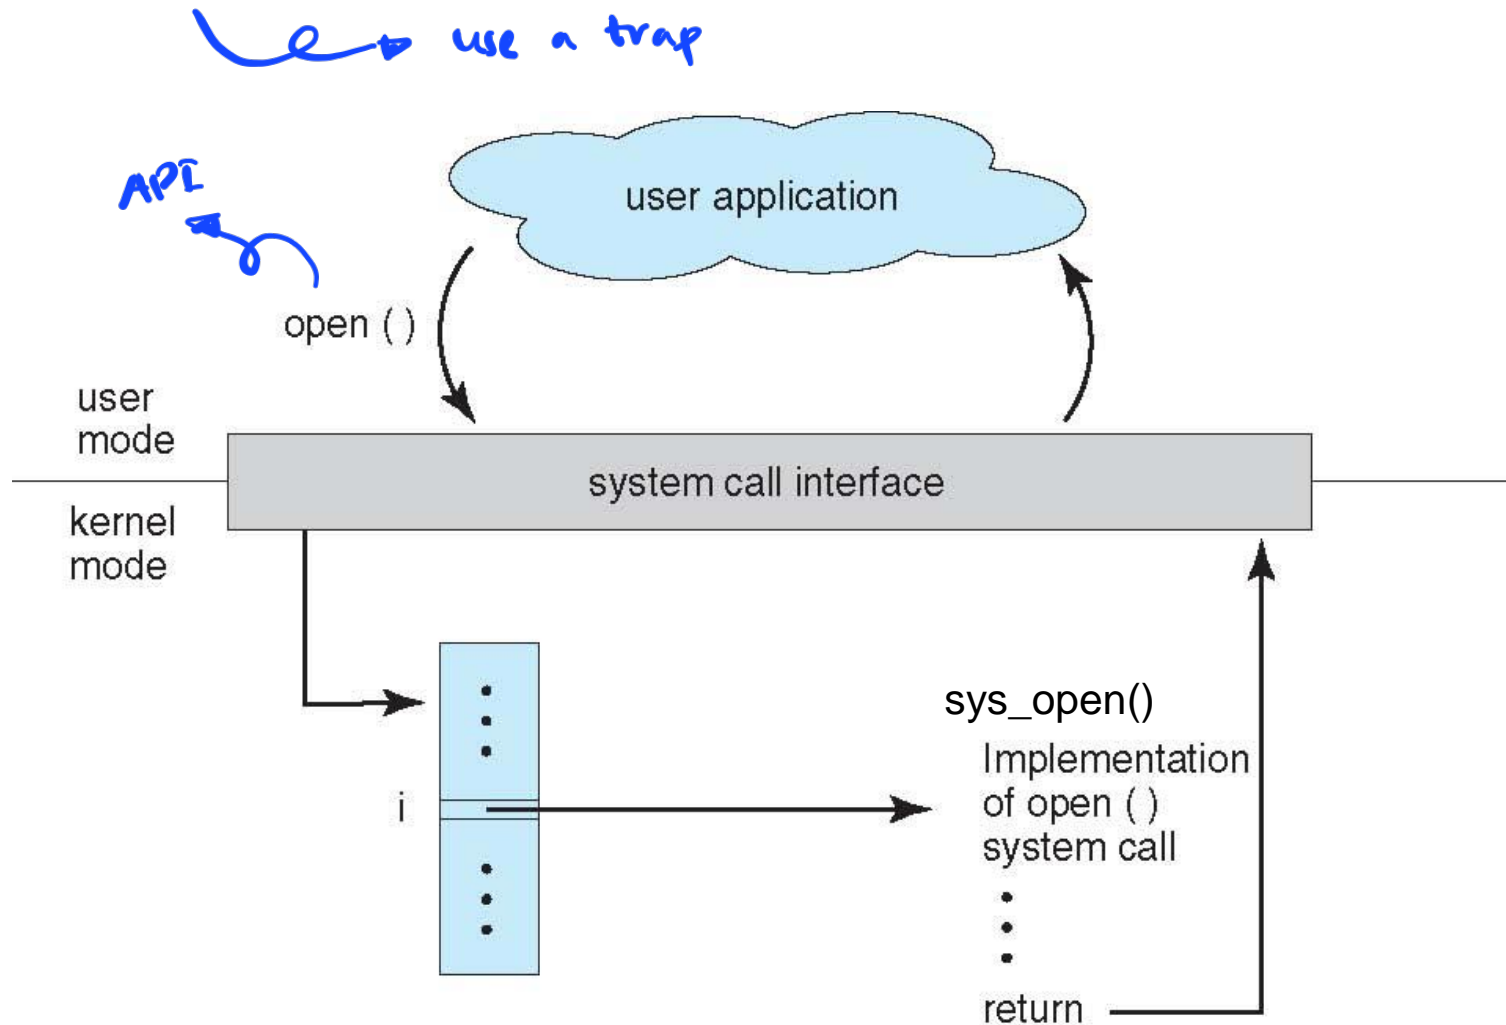
- Typically, a number associated with each system call
  - **System-call interface** maintains a table indexed according to these numbers
- The system call interface invokes intended system call in OS kernel and returns status of the system call and any return values
- The caller need know nothing about how the system call is implemented
  - Just needs to **obey API** and **understand what OS will do as a result call**
  - Most details of OS interface hidden from programmer by API
    - ▶ Managed by run-time support library (set of functions built into libraries included with compiler)







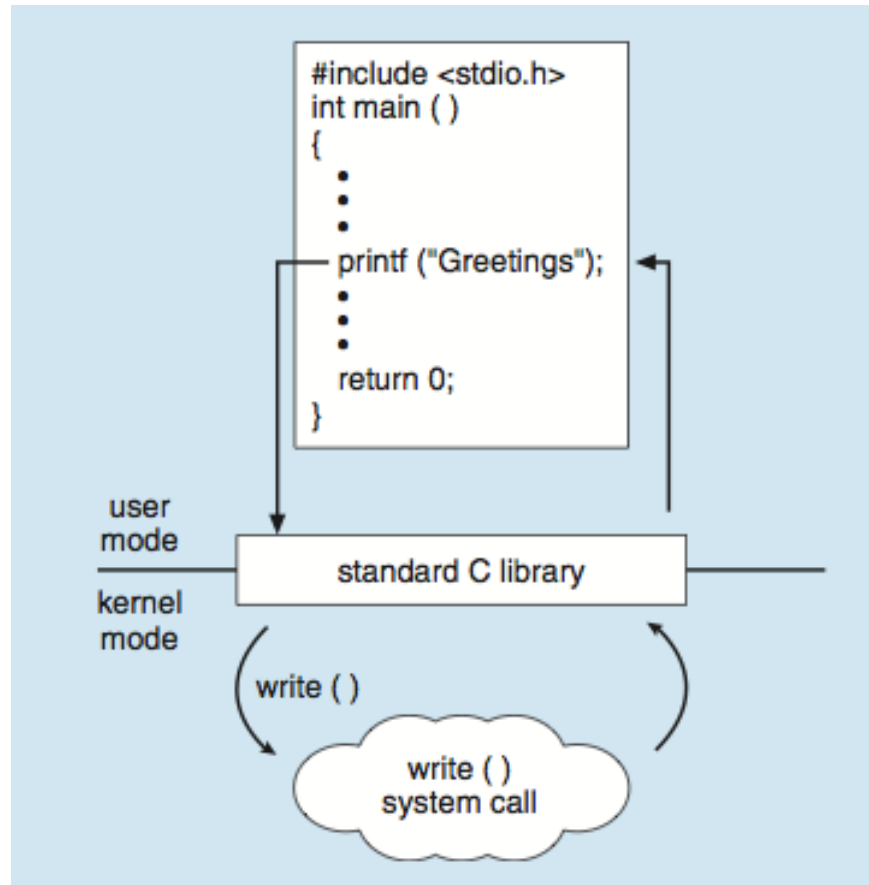
# API – System Call – OS Relationship





# Standard C Library Example

- C program invoking printf() library call, which calls write() system call



# Quizzes

- ❑ An operating system can have more than one command interpreter which a user can choose. ✓
- ❑ Every operating system has a graphical user interface. ✗
- ❑ In the Linux shell, it is possible to concatenate commands, so that the output of a command becomes the input of another command. ✓
- ❑ Reading user-input from the keyboard requires system calls. ✓  
instead, it's in the **KERNEL**
- ❑ Typically, every user process maintains a table of system calls. ✗  
which is in the OS
- ❑ To use system calls correctly in a program, one needs to know how the systems calls are implemented. ✗  
instead, use **API**



# Types of System Calls

- Process control
  - **create process, terminate process** (i.e., explicitly ending a specific process) API: Fork
  - **end** (i.e., normal end of execution), **abort** (i.e., exceptions cause OS intervention) API: exit
  - **load** ("turn a program into a process": create the memory image of the process), **execute** (CPU dispatching)
  - **get process attributes, set process attributes** (e.g., process identifier, priority)
  - **wait for time, wait event** (e.g., termination of a child process, data ready to be read in a pipe), **signal event**
  - **allocate** and **free** memory
- Issues:
  - ▶ Dump memory if error
  - ▶ **Debugger** for determining **bugs, single step** execution
  - ▶ Background/foreground execution
  - ▶ **Locks** for managing access to shared data between processes

raises  
an error





# Examples of Windows and Unix System Calls

	Windows	Unix
Process Control	CreateProcess() ExitProcess() WaitForSingleObject()	fork() exit() wait()
File Manipulation	CreateFile() ReadFile() WriteFile() CloseHandle()	open() read() write() close()
Device Manipulation	SetConsoleMode() ReadConsole() WriteConsole()	ioctl() read() write()
Information Maintenance	GetCurrentProcessID() SetTimer() Sleep()	getpid() alarm() sleep()
Communication	CreatePipe() CreateFileMapping() MapViewOfFile()	pipe() shmget() mmap()
Protection	SetFileSecurity() InitializeSecurityDescriptor() SetSecurityDescriptorGroup()	chmod() umask() chown()

UNX  
treats  
I/O devices  
like files





# Types of System Calls

- File management
  - create file, delete file
  - open, close file
  - read, write, reposition
  - get and set file attributes
  
- Device management
  - request device, release device
    - ▶ Exclusive use. **Deadlock** → two processes want to access the same resource
  - read, write, reposition
    - ▶ Physical vs. abstract devices. Similarity between files and devices
  - get device attributes, set device attributes
  - logically attach or detach devices





# Types of System Calls

---

- Information maintenance
  - get time or date, set time or date
  - get system data, set system data
  - get and set process, file, or device attributes
  
- Communications
  - create, delete communication connection
    - ▶ open, close, accept, wait
  - send, receive messages if **message passing model** to **host name** or **process name**
    - ▶ From **client** to **server**
  - **Shared-memory model** create and gain access to memory regions
    - ▶ Create, attach





# Types of System Calls

---

- Protection (control access to resources)
  - Get and set permissions
  - Allow and deny user access







# System Programs (Utilities)

*use kernel services as an intermediate*

- **System programs** provide a convenient environment for program development and execution. They can be divided into:
  - File manipulation (e.g., file manager)
  - Status information (e.g., task manager)
  - File editing (graphical or command-line e.g., vi) & content search (e.g., find, which,...)
  - Programming-language support (JVM, gcc,...)
  - Program loading and execution
  - Communications (e.g., web browsers)
  - Background services (daemons)
  - Application programs
- Most users' view of the operation system is defined by system programs, **not the actual system calls**
- Utilities use kernel services but are not part of it





# Operating System Structure

---


- General-purpose OS is very large program
- Various ways to structure one as follows
  - (knowing all details is less and less important nowadays)





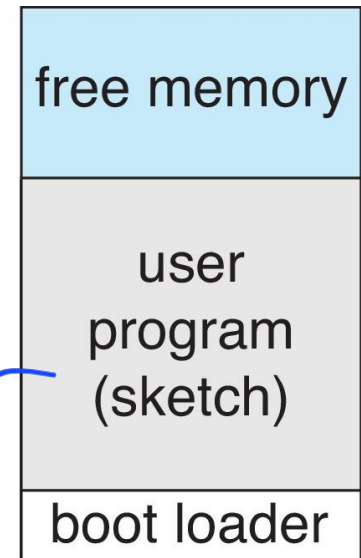
# Example: Arduino

- Single-tasking *only one thing at a time*
- No operating system
- Programs (sketch) loaded via USB into flash memory
- Single memory space
- Boot loader loads program
- Program exit -> shell reloaded



(a)

At system startup



(b)

running a program

*uploaded by USB*






# Example: MS-DOS

- Single-tasking
- Interpreter launched at start-up
- Does not create new process.
  - Overwrites the interpreter in the memory
- Output of the program saved in memory
- At the end, the interpreter is reloaded from disk into memory
- Output code made available



(a)

At system startup



(b)

running a program






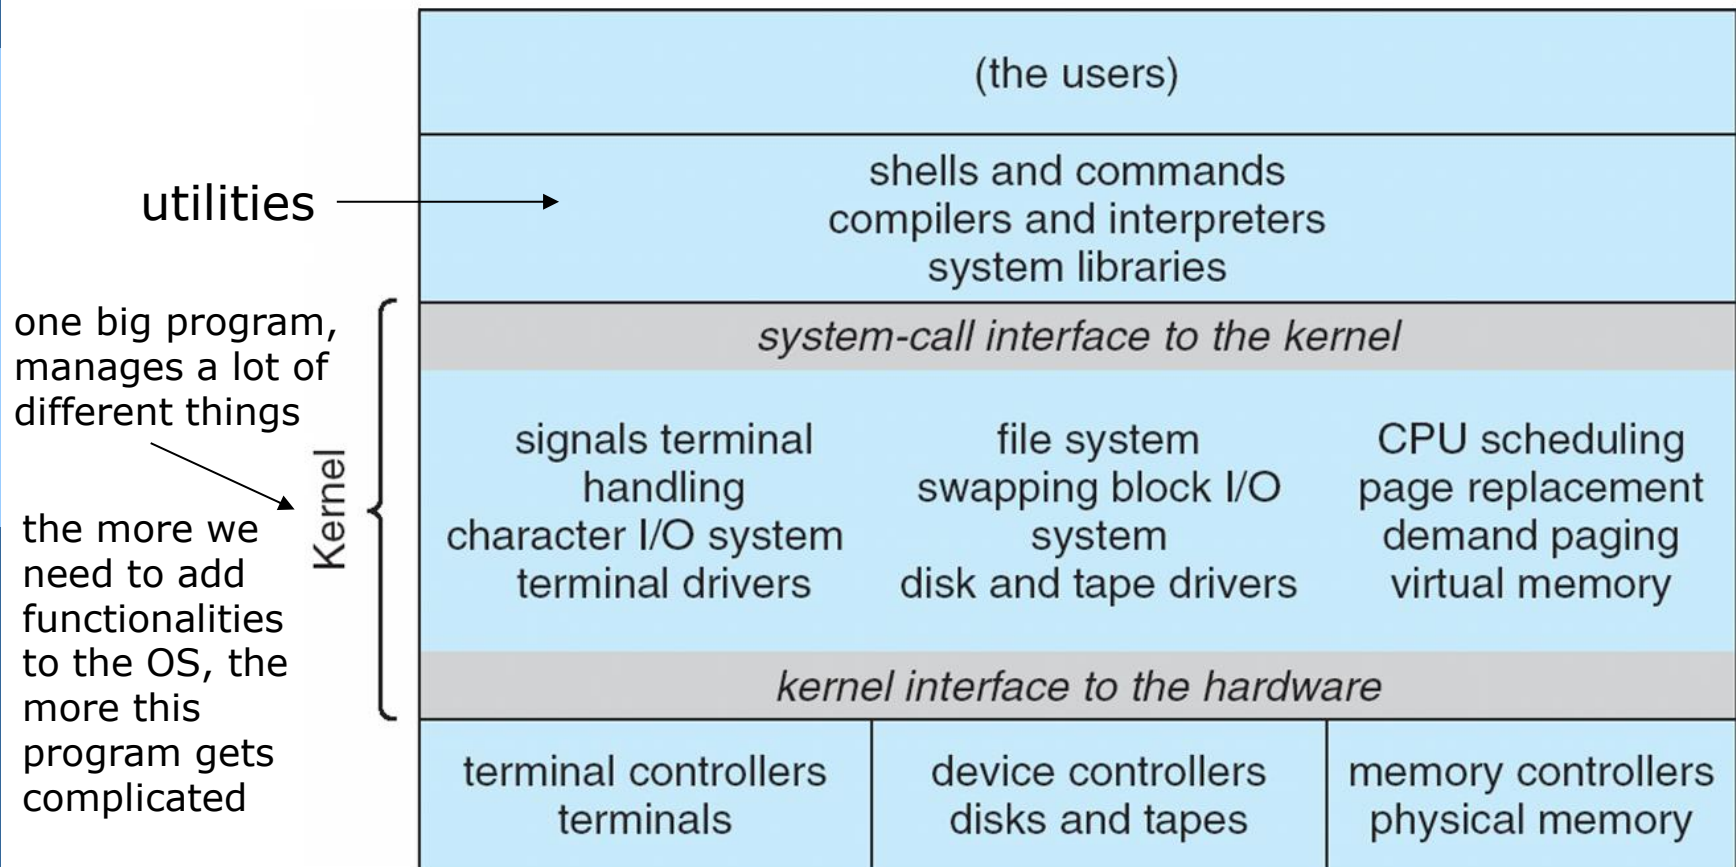
# UNIX → layers

- UNIX – limited by hardware functionality, the original UNIX operating system had limited structuring. The UNIX OS consists of two separable parts
  - Systems programs
  - The kernel
    - ▶ Consists of everything below the system-call interface and above the physical hardware
    - ▶ Provides the file system, CPU scheduling, memory management, and other operating-system functions; a large number of functions for one level





# Traditional UNIX System Structure

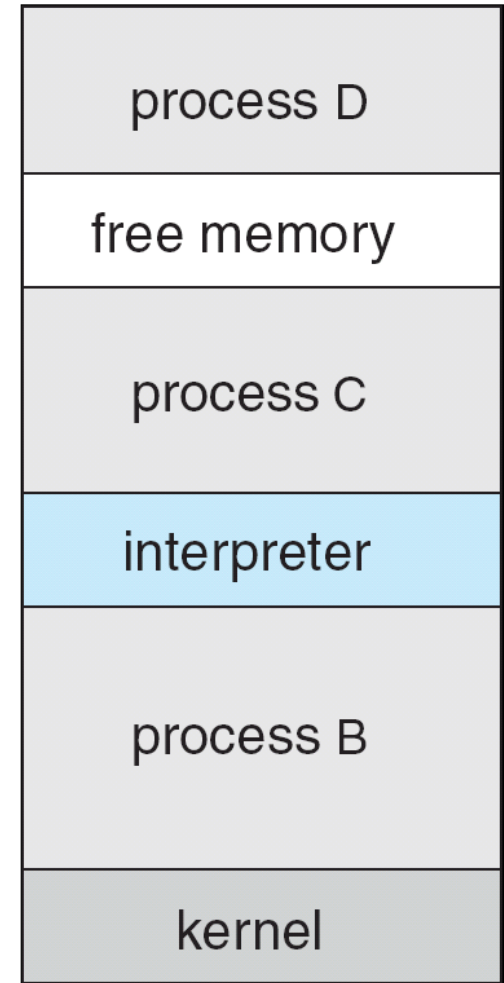




# Example: FreeBSD

- Unix variant
- Multitasking
- User login -> invoke user's choice of shell
- Shell executes `fork()` system call to create process
  - Executes `exec()` to load program into process
  - Shell waits for process to terminate or continues with user commands
- Process exits with
  - code of 0 – no error, or
  - > 0 – error code

remains  
in the  
memory





# Modules

- Most modern operating systems (e.g., VxWorks) implement loadable kernel modules → more "efficient" kernel that only loads "necessary" modules to the kernel
  - Uses object-oriented approach
  - Each core component is separate
  - Each talks to the others over known interfaces
  - Each is loadable as needed within the kernel





# Quizzes

- What **system calls** must be executed by a shell in order to **start a new process**? *fork()*
- **[T]** [F] The operating system is always **stored** in the hard disk *(storage device)*
- [T] **[F]** Utilities are loadable kernel modules

*↳ additional functions  
to interact w/ the kernel*